

## AI WORK / INNOVATION

# Root-cause analysis

IDEATION-2026-09-06-AI-WORK-01

IDEATION v0.2 | 6 September 2026

Evidence-led exploration of the Personal AI Work Coach / AI Workflow Transformation opportunity.

Research complete. Commercial evidence gates remain conditional.

# Step 3 - Root-cause analysis

Run: IDEATION-2026-09-06-AI-WORK-01. Version: 1.0. Method: IDEATION v0.2. Owner: Nadav / Codex. Date: 2026-09-06. Status: final desk-research artifact. Boundary: professional AI workflow improvement; no deployment or customer testing. Inputs: source handoff, source register and prior steps in this run.

## Fishbone breadth

People: limited attention, unclear operator, mismatch between sponsor and user. Process: undocumented variations, no acceptance standard, no dedicated setup time. Tools: account permissions, fragmented context, brittle integrations. Measurement: missing baseline, selective examples, excluded review time. Incentives: fear of disruption, no reward for saved capacity, provider rewarded for installation. Environment: sensitive information, changing policies and capabilities.

## Causal chains and alternatives

P02 -> proposed process not used -> setup/ownership handoff unfinished -> account access and a scheduled working session absent -> delivery defined as documents instead of repeatable operation. RC01: ownership/access coordination gap. EV03 supports one case (E1). Alternatives: low pain, insufficient motivation, wrong workflow, or scheduling unrelated to the design.

P03 -> improvement unknown -> unlike executions compared -> no baseline/rubric -> outcome measurement not part of completion. RC02: no agreed acceptance-and-cost contract. EV03 and EV06 support investigation, not universal causality.

P01 -> attractive task chosen over valuable task -> visible time savings dominate -> frequency/review cost not represented -> no comparable opportunity accounting. RC03: incomplete economic representation (EV01, E0). Alternative: buyer already knows the right task but cannot obtain access.

P04 -> exceptions go unresolved -> agent output treated as finished -> decision rights and recovery procedure unclear. RC04: absent exception ownership (E0). Alternative: inherently unreliable process or poor data.

P05 -> updates ignored or over-adopted -> capability announcements do not map to local tests -> no maintained representative examples/version record. RC05: missing link between external changes and local proof (E0). Alternative: current process is already adequate; no recurring job exists.

## Evaluation without composite ranking

Contribution / Evidence strength / Addressability, each 1-5 high good: RC01 4/2/4; RC02 5/3/4; RC03 3/1/3; RC04 4/1/3; RC05 3/1/2. Contribution is E0 judgment. RC01/RC02 carry into primary exploration; RC03 supports intake; RC04/RC05 remain conditional alternatives. No composite root-cause ranking, so sensitivity is not applicable.

## Red Team and gate

Challenge: missing measurement does not prove there was value to measure. Response: separate evidence infrastructure from benefit; explicitly permit no-change decisions. Challenge: no user interviews means a Five Whys chain can be fiction. Response: label RC03-RC05 E0, carry them only as hypotheses, and forbid a 'validated cause' claim. CONDITIONAL PASS into divergence. Outputs RC01-RC05 linked to P01-P05; EV03/EV06. No causal hypothesis is proven across customers.

## Source and evidence register

All sources accessed 2026-09-06. Grades apply only to the stated claim, never automatically to this business. E0 = hypothesis; E1 = anecdote/proxy; E2 = working support; E3 = direct primary evidence; E4 = repeated convergent primary evidence. No E4 demand evidence exists.

- SRC01 / EV01: User's Innovation handoff, preserved as source-handoff.txt. E3 for the requested research boundary; E0 for its unverified market and demand assertions. Direction: mixed. Earlier Storm/Red Team full transcripts were not reviewed; their conclusions are represented by this handoff.

- SRC02 / EV02: Local docs/service/offer.md and delivery-workflow.md, dated 2026-08-30. E3 for documented offer design, unknown for willingness to pay. Direction: supports an existing service vehicle. Existing AI Value Sprint: 7-10 business days, one Quick Win implemented.

- SRC03 / EV03: Local docs/clients/2026-08-30-michal-engagement-retro.md. E1 for generalizable delivery learning; E3 for the record's stated limitations. A related-party pilot reached setup, but captured no baseline and does not establish arm's-length demand. Direction: mixed. We do not repeat its claim that it proves the method works as a validated outcome.
- SRC04 / EV04: Local docs/gtm/fit-call-log.md, read today: zero recorded fit calls, qualified workflows, proposal requests and paid engagements. E3 for the record, not proof that no unlogged conversations occurred. Direction: contradicts any claim of commercial validation. Existing review: 2026-09-11 or ten calls.
- SRC05 / EV05: NBER, Generative AI at Work, <https://www.nber.org/papers/w31161> . Primary research surfaced in search; page fetch returned 403. E2 in this run because only the abstract/search evidence was accessible. Direction: supports possible task-specific gains, not universal uplift or consultant demand. Avoid transferring study effect sizes to this offer.
- SRC06 / EV06: METR, 2026-02-24 study-design update, <https://metr.org/blog/2026-02-24-uplift-update/> . E3 for its reported measurement problems; external validity to professional services remains uncertain. Direction: mixed. Newer study selection and timing problems prevent a clean current uplift estimate; the early-2025 slowdown cannot establish current model performance.
- SRC07 / EV07: Anthropic, Building effective agents, <https://www.anthropic.com/engineering/building-effective-agents> . E3 for vendor architecture guidance, E1 for commercial inference. Direction: supports using simple designs and increasing complexity only when justified. Not independent validation of this product.
- SRC08 / EV08: Zapier support and partner documentation, <https://help.zapier.com/hc/en-us/articles/8496213764877-Get-help-and-support-with-Zapier> and <https://docs.zapier.com/integrations/quickstart/app-developer-services> . E3 for availability of implementation/maintenance partners. Direction: contradicts claims that workflow implementation is empty competitive space. No verified competitor prices or small-firm outcomes.
- SRC09 / EV09: Celonis Enterprise AI, <https://www.celonis.com/solutions/ai> . E3 for advertised process intelligence/digital-twin positioning; E0 for efficacy in our target segment. Direction: contradicts novelty of a broad work digital twin. Does not establish a direct small-firm substitute.
- SRC10 / EV10: Microsoft Work Trend Index, <https://www.microsoft.com/en-us/worklab/work-trend-index> . E3 for vendor research agenda; E1 for this opportunity. Direction: supports organizational adaptation as a live topic. Vendor incentives and survey selection apply; not willingness-to-pay evidence.
- SRC11 / EV11: Local docs/service/value-proposition.md. E3 for recorded positioning and founder-background claims, not independent verification. Direction: supports provisional fit with risk-aware professional services; reachable buyers and delivery capacity remain unverified.

## Shared assumptions, risks, and contradictions

ASM01: reachable buyers have a costly recurring workflow (E0). ASM02: they will pay for implementation plus proof (E0). ASM03: cases share reusable structure (E0). ASM04: safe representative samples can be obtained (E0). ASM05: customers will later pay for revalidation (E0).

RISK01: verification/setup overhead erases gains. RISK02: no paying buyer. RISK03: bespoke delivery consumes founder capacity. RISK04: sensitive context leaks even after names are removed. RISK05: bundled platform features absorb the software. RISK06: provider grades its own work too generously.

Contradictions: general productivity opportunity (EV05/EV10) coexists with local absence of measured outcomes (EV03/EV04); these concern different populations and require a measured local pilot. More capable agents (scenario, not forecast) can increase delegation while reducing the price of setup; survival of independent assurance remains ASM05. Privacy differentiation in local marketing (EV11) is not proof competitors lack it (EV08).