

## **Divergent ideation**

IDEATION-2026-09-06-AI-WORK-01

IDEATION v0.2 | 6 September 2026

Evidence-led exploration of the Personal AI Work Coach / AI Workflow Transformation opportunity.

Research complete. Commercial evidence gates remain conditional.

# Step 4 - Divergent ideation

Run: IDEATION-2026-09-06-AI-WORK-01. Version: 1.0. Method: IDEATION v0.2. Owner: Nadav / Codex. Date: 2026-09-06. Status: final desk-research artifact. Boundary: professional AI workflow improvement; no deployment or customer testing. Inputs: source handoff, source register and prior steps in this run.

## Intervention space before selection

SF01 implementation/coordination (RC01): reserve live setup time; choose workflow owner; use existing tools; remove a handoff; standardize inputs; eliminate the task entirely.

SF02 measurement/assurance (RC02): define accepted output; count review time; blind comparisons; build replay cases; audit failed changes; revalidate vendor deliverables.

SF03 prioritization (RC03): rank by frequency and avoidable cost; audit bottlenecks; compare buying existing software; outsource a result; refuse low-value automation; redeploy saved capacity.

SF04 delegation/governance (RC04): exception queue; decision-right contract; shadow agent; deterministic checks; rollback rehearsal; manual fallback.

SF05 adaptive optimization (RC05): event-triggered retest; version ledger; capability-to-case matching; model migration; drift alarm; quarterly opt-in review.

SF06 capability transfer (RC01/RC04): apprenticeship on real task; observed independent run; supervision practice; failure rehearsal; portable method; no new product.

These 36 primitives preserve materially different interventions; the 14 concepts below are combinations, not the limit of divergence. Card family tags are provisional; Step 5 evaluates mechanisms rather than accepting every card element.

## Analogies and limits

Personal trainer: observe performance; reject mandatory lesson schedules. Physical therapist: regain independence from provider, not from useful tools. Management consultant: prioritize economics; risk slide-only delivery. Automation consultant: own setup; risk automating the wrong process. Apprenticeship: learn on useful work; risk expert bottleneck. Security scanner: event-driven discovery; risk alert fatigue. GPS: intervene at the moment of a decision; risk inaccessible work context. IDE: use actual artifacts and verification; business outcomes often lack executable tests. DevOps: version, monitor, rollback; quality can be subjective. Financial advisor: allocate scarce improvement effort; avoid unverifiable 'frontier returns'. Performance coach: measure outcomes; provider is not an independent judge. Operating system: carry context; integration costs explode. Digital twin: simulate before change; behavioral reality may not replay. Additional analogies: building commissioning (acceptance before occupancy), metrology lab (calibrated independent measurement), and industrial maintenance (repair only when economically justified).

### ***101 - Workflow Acceptance Sprint (service; vertical)***

Implement one recurring professional-services workflow and demonstrate accepted outputs.

Customer: Practice owner. Trigger: Weekly reporting backlog. Core job: Get one process running with credible proof. Product object: Accepted workflow version.

Interaction: Observe, configure, replay, hand over. Human/AI split: Founder implements; AI drafts; client approves.

Why pay: Operational setup and accountable completion. Why ChatGPT/Claude alone may be insufficient: Assistant may design it; buyer still needs execution/access/acceptance ownership. Revenue: Fixed scoped fee. Retention: Next painful workflow, separately purchased.

Technical complexity: Low-medium. Founder fit: High provisional: existing practice. Biggest failure mode: Bespoke support consumes margin. All demand/fit ratings are E0 planning hypotheses; links: SF01/SF02/SF06.

### ***102 - Fractional Work Systems Partner (service)***

Manage a firm's recurring improvement backlog.

Customer: Small-firm owner. Trigger: Several workflows breaking or changing. Core job: Allocate and deliver improvements. Product object: Improvement portfolio.

Interaction: Monthly review plus scoped implementation. Human/AI split: Human prioritizes and owns change; AI analyzes.

Why pay: A responsible operator. Why ChatGPT/Claude alone may be insufficient: Coordination and delivery capacity are purchased. Revenue: Retainer with capacity cap. Retention: Actual incidents and new work.

Technical complexity: Medium. Founder fit: Medium: recurring time commitment. Biggest failure mode: Retainer becomes unlimited support. All demand/fit ratings are E0 planning hypotheses; links: SF01/SF03/SF05.

### ***103 - Independent Workflow Acceptance Lab (service; contrarian)***

Test AI workflows built by the buyer or another vendor.

Customer: Buyer with an existing implementation. Trigger: Renewal, handover or disputed ROI. Core job: Decide accept, fix or reject. Product object: Acceptance report.

Interaction: Redacted samples and a predefined test protocol. Human/AI split: Evaluator owns rubric; AI runs checks.

Why pay: Independent decision evidence. Why ChatGPT/Claude alone may be insufficient: Separation from implementer and reproducible assessment. Revenue: Per assessment; revalidation optional. Retention: Version change or renewal.

Technical complexity: Low-medium. Founder fit: High technical fit; buyer access unproven. Biggest failure mode: Buyer expects testing free from implementer. All demand/fit ratings are E0 planning hypotheses; links: SF02.

### ***104 - Work Upgrade Ledger (software)***

Track workflow versions, outcomes and ownership across existing tools.

Customer: Small team workflow owner. Trigger: Second workflow or tool change. Core job: Know what works and why. Product object: Workflow version plus evidence.

Interaction: Import records; compare versions; approve change. Human/AI split: Software records; humans judge ambiguous quality.

Why pay: Reduced loss of operational knowledge. Why ChatGPT/Claude alone may be insufficient: Durable structured records; platforms can still copy this. Revenue: Per team subscription. Retention: Run evidence and change reviews.

Technical complexity: Medium. Founder fit: Medium. Biggest failure mode: Data entry outweighs value. All demand/fit ratings are E0 planning hypotheses; links: SF02/SF05.

### ***105 - AI Opportunity Scanner (software)***

Propose retests when a tool change might improve a known workflow.

Customer: AI-forward operator. Trigger: Relevant capability or price change. Core job: Find worthwhile upgrades. Product object: Candidate retest.

Interaction: Workflow registry plus targeted alerts. Human/AI split: AI maps changes; human approves tests.

Why pay: Avoid monitoring every release. Why ChatGPT/Claude alone may be insufficient: Local history plus actual tests; generic recommendations are insufficient. Revenue: Subscription. Retention: Only material local opportunities.

Technical complexity: High. Founder fit: Low before a dataset. Biggest failure mode: False positives and switching cost. All demand/fit ratings are E0 planning hypotheses; links: SF05.

### ***106 - Workflow Simulator (software)***

Replay proposed process variants against representative historical cases.

Customer: Operations lead. Trigger: Before deploying a redesign. Core job: Estimate quality and cost tradeoffs. Product object: Replayable case set.

Interaction: Sandbox compare view. Human/AI split: AI executes variants; humans adjudicate.

Why pay: Avoid failed process changes. Why ChatGPT/Claude alone may be insufficient: Real case corpus, connectors and reproducible replay. Revenue: Usage plus workspace fee. Retention: Every material revision.

Technical complexity: High. Founder fit: Medium technical; low data fit. Biggest failure mode: Replay does not capture real-world behavior. All demand/fit ratings are E0 planning hypotheses; links: SF02/SF05.

### ***107 - Delegation Control Desk (software; agent-dominant)***

Route agent exceptions and approvals to accountable humans.

Customer: Team supervising agents. Trigger: Multiple agents produce unresolved cases. Core job: Resolve exceptions without inspecting everything. Product object: Decision right and exception.

Interaction: Queue with escalation and fallback. Human/AI split: Agents execute; humans handle exceptions.

Why pay: Reliable accountability and response. Why ChatGPT/Claude alone may be insufficient: Cross-system routing and policy; native platforms may absorb it. Revenue: Per active workflow/team. Retention: Real exceptions.

Technical complexity: High. Founder fit: Medium. Biggest failure mode: Exception volume makes supervision expensive. All demand/fit ratings are E0 planning hypotheses; links: SF04.

### ***108 - Agent Migration Clinic (service; agent-dominant)***

Move an established process to a new execution model only after acceptance tests.

Customer: Owner of a brittle automation. Trigger: Provider change or capability step. Core job: Migrate without losing reliability. Product object: Verified migration.

Interaction: Shadow run, cutover, rollback. Human/AI split: Founder migrates; AI executes; client approves.

Why pay: Controlled continuity. Why ChatGPT/Claude alone may be insufficient: Actual integration, comparison and rollback responsibility. Revenue: Fixed migration fee. Retention: Migrations and failures.

Technical complexity: Medium-high. Founder fit: Medium-high. Biggest failure mode: Migrations too infrequent or bespoke. All demand/fit ratings are E0 planning hypotheses; links: SF04/SF05.

### ***109 - Consultant Evidence Brief Pipeline (service; vertical)***

Install a repeatable source-linked client briefing process.

Customer: Boutique consulting lead. Trigger: Repeated weekly client briefs. Core job: Deliver accepted research briefs with less total effort. Product object: Evidence-backed brief.

Interaction: Approved inputs to draft to reviewer. Human/AI split: AI extracts/drafts; consultant verifies conclusions.

Why pay: Capacity and consistent traceability. Why ChatGPT/Claude alone may be insufficient: Workflow-specific setup and acceptance; native research is a strong substitute. Revenue: Implementation fee. Retention: New briefing types/clients.

Technical complexity: Medium. Founder fit: High provisional. Biggest failure mode: Review costs erase drafting gains. All demand/fit ratings are E0 planning hypotheses; links: SF01/SF02.

### ***110 - Recruiter Intake-to-Shortlist System (software; vertical)***

Structure approved role criteria and evidence for human-reviewed shortlists.

Customer: Recruiting firm lead. Trigger: Repeated intake/shortlist backlog. Core job: Reduce coordination and inconsistent evidence. Product object: Candidate evidence packet.

Interaction: Intake form to review queue. Human/AI split: AI organizes; recruiter decides.

Why pay: Faster consistent intake. Why ChatGPT/Claude alone may be insufficient: Team workflow and evidence; ATS incumbents may bundle. Revenue: Team subscription. Retention: Each hiring assignment.

Technical complexity: High. Founder fit: Low without domain/data access. Biggest failure mode: Sensitive data and unfair ranking. All demand/fit ratings are E0 planning hypotheses; links: SF01/SF04.

### ***111 - Work Upgrade Protocol (methodology)***

A portable process people run inside their existing assistant.

Customer: Capable DIY professional. Trigger: A specific recurring task hurts. Core job: Redesign and measure work independently. Product object: Reusable protocol.

Interaction: Guided worksheet inside current tools. Human/AI split: User performs; AI guides.

Why pay: Convenient expert structure, if anything. Why ChatGPT/Claude alone may be insufficient: Often assistant alone is enough; no strong moat claimed. Revenue: Workshop/license or free acquisition tool. Retention: Occasional reuse.

Technical complexity: Low. Founder fit: High. Biggest failure mode: Free alternatives remove willingness to pay. All demand/fit ratings are E0 planning hypotheses; links: SF06.

### ***I12 - No New Product (no standalone business)***

Use the method internally in NY AI Solutions and buy existing software where adequate.

Customer: Existing consulting prospect. Trigger: Current process needs fixing. Core job: Choose the simplest viable intervention. Product object: Client outcome.

Interaction: Existing fit call and sprint. Human/AI split: Current operator and tools.

Why pay: Only existing service delivery earns revenue. Why ChatGPT/Claude alone may be insufficient: No claim that another product is needed. Revenue: Existing practice revenue. Retention: Existing client needs.

Technical complexity: Low. Founder fit: Very high. Biggest failure mode: Never testing a repeatable product opportunity. All demand/fit ratings are E0 planning hypotheses; links: SF03/SF06.

### ***I13 - Capacity Redeployment Sprint (service; vertical)***

Turn saved time into a specific additional client-serving activity.

Customer: Agency/practice owner. Trigger: AI saves time but revenue or backlog does not improve. Core job: Convert capacity into useful output. Product object: Accepted capacity allocation.

Interaction: Baseline bottleneck; change process; measure downstream use. Human/AI split: Human changes operating policy; AI reduces tasks.

Why pay: Real throughput rather than theoretical hours. Why ChatGPT/Claude alone may be insufficient: Incentives, staffing and workflow coordination. Revenue: Fixed service fee. Retention: New bottleneck.

Technical complexity: Low-medium. Founder fit: Medium. Biggest failure mode: Bottleneck lies in sales or demand, not delivery. All demand/fit ratings are E0 planning hypotheses; links: SF03.

### ***I14 - Workflow Evidence Exchange (network)***

License permissioned benchmark cases and vetted implementation patterns.

Customer: Specialist implementer. Trigger: New workflow needs trusted examples. Core job: Reduce duplicated validation effort. Product object: Reusable test case.

Interaction: Search, replay, contribute. Human/AI split: Experts curate; AI adapts.

Why pay: Faster credible evaluation. Why ChatGPT/Claude alone may be insufficient: Permissioned corpus and quality controls, if achieved. Revenue: Membership/license. Retention: New implementations.

Technical complexity: High. Founder fit: Low before repeated demand. Biggest failure mode: Cold start and nonportable private data. All demand/fit ratings are E0 planning hypotheses; links: SF02/SF05.

## **Combinations and preserved rejections**

Combine audit + implementation + acceptance + operator handoff into I01; scanner stays optional. Combine ledger + replay + conditional retesting into I04/I06, without a full work digital twin. Combine consultant vertical + representative examples + implementation into I09. Combine independent evaluation + migration into I03/I08 only when buyer commissions both and evaluator independence is not misrepresented.

Reject automatic optimization toward maximum autonomy: it optimizes a proxy. Defer universal screen observation: disproportionate privacy and integration burden. Preserve I12 as a valid outcome, not a straw man. No concept is selected at this stage.

## **Gate**

PASS for breadth. Inputs RC01-RC05; outputs I01-I14, SF01-SF06. Evidence EV01/EV07-EV11 informs analogies and competition; commercial appeal remains E0. No composite score in this step. Next: mechanism evidence gate.

## Source and evidence register

All sources accessed 2026-09-06. Grades apply only to the stated claim, never automatically to this business. E0 = hypothesis; E1 = anecdote/proxy; E2 = working support; E3 = direct primary evidence; E4 = repeated convergent primary evidence. No E4 demand evidence exists.

- SRC01 / EV01: User's Innovation handoff, preserved as source-handoff.txt. E3 for the requested research boundary; E0 for its unverified market and demand assertions. Direction: mixed. Earlier Storm/Red Team full transcripts were not reviewed; their conclusions are represented by this handoff.

- SRC02 / EV02: Local docs/service/offer.md and delivery-workflow.md, dated 2026-08-30. E3 for documented offer design, unknown for willingness to pay. Direction: supports an existing service vehicle. Existing AI Value Sprint: 7-10 business days, one Quick Win implemented.

- SRC03 / EV03: Local docs/clients/2026-08-30-michal-engagement-retro.md. E1 for generalizable delivery learning; E3 for the record's stated limitations. A related-party pilot reached setup, but captured no baseline and does not establish arm's-length demand. Direction: mixed. We do not repeat its claim that it proves the method works as a validated outcome.

- SRC04 / EV04: Local docs/gtm/fit-call-log.md, read today: zero recorded fit calls, qualified workflows, proposal requests and paid engagements. E3 for the record, not proof that no unlogged conversations occurred. Direction: contradicts any claim of commercial validation. Existing review: 2026-09-11 or ten calls.

- SRC05 / EV05: NBER, Generative AI at Work, <https://www.nber.org/papers/w31161> . Primary research surfaced in search; page fetch returned 403. E2 in this run because only the abstract/search evidence was accessible. Direction: supports possible task-specific gains, not universal uplift or consultant demand. Avoid transferring study effect sizes to this offer.

- SRC06 / EV06: METR, 2026-02-24 study-design update, <https://metr.org/blog/2026-02-24-uplift-update/> . E3 for its reported measurement problems; external validity to professional services remains uncertain. Direction: mixed. Newer study selection and timing problems prevent a clean current uplift estimate; the early-2025 slowdown cannot establish current model performance.

- SRC07 / EV07: Anthropic, Building effective agents, <https://www.anthropic.com/engineering/building-effective-agents> . E3 for vendor architecture guidance, E1 for commercial inference. Direction: supports using simple designs and increasing complexity only when justified. Not independent validation of this product.

- SRC08 / EV08: Zapier support and partner documentation, <https://help.zapier.com/hc/en-us/articles/8496213764877-Get-help-and-support-with-Zapier> and <https://docs.zapier.com/integrations/quickstart/app-developer-services> . E3 for availability of implementation/maintenance partners. Direction: contradicts claims that workflow implementation is empty competitive space. No verified competitor prices or small-firm outcomes.

- SRC09 / EV09: Celonis Enterprise AI, <https://www.celonis.com/solutions/ai> . E3 for advertised process intelligence/digital-twin positioning; E0 for efficacy in our target segment. Direction: contradicts novelty of a broad work digital twin. Does not establish a direct small-firm substitute.

- SRC10 / EV10: Microsoft Work Trend Index, <https://www.microsoft.com/en-us/worklab/work-trend-index> . E3 for vendor research agenda; E1 for this opportunity. Direction: supports organizational adaptation as a live topic. Vendor incentives and survey selection apply; not willingness-to-pay evidence.

- SRC11 / EV11: Local docs/service/value-proposition.md. E3 for recorded positioning and founder-background claims, not independent verification. Direction: supports provisional fit with risk-aware professional services; reachable buyers and delivery capacity remain unverified.

## Shared assumptions, risks, and contradictions

ASM01: reachable buyers have a costly recurring workflow (E0). ASM02: they will pay for implementation plus proof (E0). ASM03: cases share reusable structure (E0). ASM04: safe representative samples can be obtained (E0). ASM05: customers will later pay for revalidation (E0).

RISK01: verification/setup overhead erases gains. RISK02: no paying buyer. RISK03: bespoke delivery consumes founder capacity. RISK04: sensitive context leaks even after names are removed. RISK05: bundled platform features absorb the software. RISK06: provider grades its own work too generously.

Contradictions: general productivity opportunity (EV05/EV10) coexists with local absence of measured outcomes (EV03/EV04); these concern different populations and require a measured local pilot. More capable agents (scenario, not forecast) can increase delegation while reducing the price of setup; survival of independent assurance remains ASM05. Privacy differentiation in local marketing (EV11) is not proof competitors lack it (EV08).