

AI WORK / INNOVATION

Solution validation and evidence gate

IDEATION-2026-09-06-AI-WORK-01

IDEATION v0.2 | 6 September 2026

Evidence-led exploration of the Personal AI Work Coach / AI Workflow Transformation opportunity.

Research complete. Commercial evidence gates remain conditional.

Step 5 - Solution validation and evidence gate

Run: IDEATION-2026-09-06-AI-WORK-01. Version: 1.0. Method: IDEATION v0.2. Owner: Nadav / Codex. Date: 2026-09-06. Status: final desk-research artifact. Boundary: professional AI workflow improvement; no deployment or customer testing. Inputs: source handoff, source register and prior steps in this run.

Mechanism ratings

Raw scale 1-5, higher favorable: problem fit / evidence strength / market whitespace / technical feasibility / product leverage / risk manageability. Evidence score is a judgment of relevance, separate from E-grade. Baseline weights 25/25/10/15/15/10%; impact/evidence 30/35/5/10/15/5%; feasibility/risk 15/20/5/25/10/25%. Confidence low for market whitespace and product leverage, medium for bounded service feasibility. A commercial core requires direct paid demand; none passes that threshold. Retentions below are permissions for conditional concept synthesis, not validated demand.

M01 live implementation and named owner (RC01/SF01): 5/3/2/5/3/4. STRONG KEEP for service exploration; EV03/EV08. Access and repetition must be tested.

M02 acceptance contract and baseline (RC02/SF02): 5/3/3/4/5/4. CORE KEEP as evidence discipline; EV03/EV06. The mechanism prevents unsupported claims; it does not prove customers buy measurement separately.

M03 economic workflow selection (RC03/SF03): 4/1/2/5/3/4. KEEP AS ENABLER; EV01/EV02. E0 causal/demand support means interview and test, not validated prioritization algorithm.

M04 operator supervision and fallback rehearsal (RC01/RC04/SF06): 4/2/2/5/3/4. KEEP AS ENABLER; EV03/EV07. No requirement for unaided task execution.

M05 portable version/evidence ledger (RC02/RC05/SF05): 4/2/2/4/4/3. KEEP AS INFRASTRUCTURE, manual record only; EV06/EV09. Software purchase hypothesis unvalidated.

M06 autonomous opportunity scan/digital twin (RC05/SF05): 3/1/1/2/5/1. DEFER; EV09 and ASM05. No dependable local data or repeat-demand evidence.

M07 exception routing and migration (RC04/SF04): 4/2/2/3/4/2. DEFER as product, retain only manual fallback within M04; EV07. Incident frequency unproven.

M08 course/capability graph (RC01/SF06): 2/1/1/5/2/4. REJECT as product core; EV01 only. Could help particular users, but no evidence it solves selected P02/P03 better.

ID	Raw Fit/Evidence/Whitespace/Feasibility/Leverage/Risk	Baseline / impact / feasibility	Ranks	Stability
M01	5/3/2/5/3/4	3.80 / 3.80 / 4.00	2/2/1	rank-stable
M02	5/3/3/4/5/4	4.05 / 4.05 / 4.00	1/1/2	rank-stable
M03	4/1/2/5/3/4	3.05 / 2.80 / 3.45	5/6/4	rank-stable
M04	4/2/2/5/3/4	3.30 / 3.15 / 3.65	3/3/3	rank-stable
M05	4/2/2/4/4/3	3.20 / 3.15 / 3.25	4/4/5	rank-stable
M06	3/1/1/2/5/1	2.25 / 2.30 / 1.95	8/7/8	rank-stable
M07	4/2/2/3/4/2	2.95 / 3.00 / 2.75	6/5/7	rank-stable
M08	2/1/1/5/2/4	2.30 / 2.00 / 3.00	7/8/6	rank-stable

Conditional solution kernel

M01 + M02 + M04, supported by M03 and manual M05. RC01/RC02 are locally plausible, not population-validated. ASM01/02/04 are load-bearing and unresolved. M06/M07 cannot enter the committed MVP. No mechanism has passed a paid-market gate. 'Validated kernel' here means researched and bounded enough to test; a stronger label would overstate evidence.

Business models ranked for the first test

Planning rank, no measured willingness to pay: 1 B productized sprint; 2 A audit plus implementation; 3 G team workflow transformation limited to one workflow; 4 H diagnostic with an explicit implementation option; 5 I portable methodology inside existing practice; 6 D fractional AI CTO; 7 F partly outcome-linked fee; 8 E monthly optimization; 9 C premium coach; 10 L embedded product; 11 J standalone SaaS; 12 K marketplace. L is an interaction/distribution choice more than a standalone revenue model. Outcome-linked fees require an agreed counterfactual and should not shift unlimited risk to the founder. No prices recommended or published.

Pilot format comparison

These are E0 judgments relative to a solo founder, not research findings. H/M/L indicate high/medium/low, favorable only where specified.

Six-week/three-workflow accelerator: pain M, activation slow, WTP unknown, learning H, implementation H, repeat demand unknown, founder effort H, scalability L, assistant differentiation M. Too many variables for the first test.

One-week/one-workflow sprint: pain H if qualified, activation fast after access, WTP unknown, learning M, implementation H, repeat demand unknown, effort M, scalability M, differentiation M. Best evidence per founder hour if recurring executions fit the window; allow 7-10 business days plus a follow-up rather than promising one calendar week.

One-day audit: pain M, activation fastest, WTP unknown, learning M, implementation L, repeat L, effort L, scalability M, differentiation L-M. Best for screening; cannot prove delivered value.

Done-for-you implementation: pain H, activation M, WTP unknown, learning L unless handoff, implementation H, repeat M hypothesis, effort H, scalability L, differentiation M-H on execution. Cap scope and include acceptance.

Monthly service: pain depends on actual backlog, activation M, WTP unknown, learning M, implementation H, repeat unproven, effort H, scalability L, differentiation M. Premature until a second paid need.

Contextual coach: pain M, activation fast only with accessible context, WTP unknown, learning H, implementation M, repeat unproven, effort medium-to-high to build, scalability H in theory, differentiation L.

Continuous scanner: pain unproven, activation slow because a workflow model is needed, WTP unknown, learning L, implementation L by itself, repeat unproven, effort H, scalability H in theory, differentiation L for generic alerts/M for verified local retests.

Red Team and gate

Challenge: independent evaluation still costs time, and a vendor may bundle it. Response: reserve I03 as a separate falsification test, not a certainty. Challenge: digital twins/process intelligence already exist. Response: remove category-novelty claims (EV09). CONDITIONAL PASS into a test-only concept brief. No build gate passed. Objects M01-M08; risks RISK01-RISK06.

Source and evidence register

All sources accessed 2026-09-06. Grades apply only to the stated claim, never automatically to this business. E0 = hypothesis; E1 = anecdote/proxy; E2 = working support; E3 = direct primary evidence; E4 = repeated convergent primary evidence. No E4 demand evidence exists.

- SRC01 / EV01: User's Innovation handoff, preserved as source-handoff.txt. E3 for the requested research boundary; E0 for its unverified market and demand assertions. Direction: mixed. Earlier Storm/Red Team full transcripts were not reviewed; their conclusions are represented by this handoff.

- SRC02 / EV02: Local docs/service/offer.md and delivery-workflow.md, dated 2026-08-30. E3 for documented offer design, unknown for willingness to pay. Direction: supports an existing service vehicle. Existing AI Value Sprint: 7-10 business days, one Quick Win implemented.

- SRC03 / EV03: Local docs/clients/2026-08-30-michal-engagement-retro.md. E1 for generalizable delivery learning; E3 for the record's stated limitations. A related-party pilot reached setup, but captured no baseline and does not establish arm's-length demand. Direction: mixed. We do not repeat its claim that it proves the method works as a validated outcome.

- SRC04 / EV04: Local docs/gtm/fit-call-log.md, read today: zero recorded fit calls, qualified workflows, proposal requests and paid engagements. E3 for the record, not proof that no unlogged conversations occurred. Direction: contradicts any claim of commercial validation. Existing review: 2026-09-11 or ten calls.

- SRC05 / EV05: NBER, Generative AI at Work, <https://www.nber.org/papers/w31161> . Primary research surfaced in search; page fetch returned 403. E2 in this run because only the abstract/search evidence was accessible. Direction: supports possible task-specific gains, not universal uplift or consultant demand. Avoid transferring study effect sizes to this offer.
- SRC06 / EV06: METR, 2026-02-24 study-design update, <https://metr.org/blog/2026-02-24-uplift-update/> . E3 for its reported measurement problems; external validity to professional services remains uncertain. Direction: mixed. Newer study selection and timing problems prevent a clean current uplift estimate; the early-2025 slowdown cannot establish current model performance.
- SRC07 / EV07: Anthropic, Building effective agents, <https://www.anthropic.com/engineering/building-effective-agents> . E3 for vendor architecture guidance, E1 for commercial inference. Direction: supports using simple designs and increasing complexity only when justified. Not independent validation of this product.
- SRC08 / EV08: Zapier support and partner documentation, <https://help.zapier.com/hc/en-us/articles/8496213764877-Get-help-and-support-with-Zapier> and <https://docs.zapier.com/integrations/quickstart/app-developer-services> . E3 for availability of implementation/maintenance partners. Direction: contradicts claims that workflow implementation is empty competitive space. No verified competitor prices or small-firm outcomes.
- SRC09 / EV09: Celonis Enterprise AI, <https://www.celonis.com/solutions/ai> . E3 for advertised process intelligence/digital-twin positioning; E0 for efficacy in our target segment. Direction: contradicts novelty of a broad work digital twin. Does not establish a direct small-firm substitute.
- SRC10 / EV10: Microsoft Work Trend Index, <https://www.microsoft.com/en-us/worklab/work-trend-index> . E3 for vendor research agenda; E1 for this opportunity. Direction: supports organizational adaptation as a live topic. Vendor incentives and survey selection apply; not willingness-to-pay evidence.
- SRC11 / EV11: Local docs/service/value-proposition.md. E3 for recorded positioning and founder-background claims, not independent verification. Direction: supports provisional fit with risk-aware professional services; reachable buyers and delivery capacity remain unverified.

Shared assumptions, risks, and contradictions

ASM01: reachable buyers have a costly recurring workflow (E0). ASM02: they will pay for implementation plus proof (E0). ASM03: cases share reusable structure (E0). ASM04: safe representative samples can be obtained (E0). ASM05: customers will later pay for revalidation (E0).

RISK01: verification/setup overhead erases gains. RISK02: no paying buyer. RISK03: bespoke delivery consumes founder capacity. RISK04: sensitive context leaks even after names are removed. RISK05: bundled platform features absorb the software. RISK06: provider grades its own work too generously.

Contradictions: general productivity opportunity (EV05/EV10) coexists with local absence of measured outcomes (EV03/EV04); these concern different populations and require a measured local pilot. More capable agents (scenario, not forecast) can increase delegation while reducing the price of setup; survival of independent assurance remains ASM05. Privacy differentiation in local marketing (EV11) is not proof competitors lack it (EV08).