

AI WORK / INNOVATION

Competing concepts, business brief and specification

IDEATION-2026-09-06-AI-WORK-01

IDEATION v0.2 | 6 September 2026

Evidence-led exploration of the Personal AI Work Coach / AI Workflow Transformation opportunity.

Research complete. Commercial evidence gates remain conditional.

Step 6 - Competing concepts, business brief and specification

Run: IDEATION-2026-09-06-AI-WORK-01. Version: 1.0. Method: IDEATION v0.2. Owner: Nadav / Codex. Date: 2026-09-06. Status: final desk-research artifact. Boundary: professional AI workflow improvement; no deployment or customer testing. Inputs: source handoff, source register and prior steps in this run.

Three final directions

PC01 / near-term business - Workflow Acceptance Sprint, within the existing AI Value Sprint. Thesis: an owner buys one working improvement with credible acceptance and a named operator. User: workflow owner and practice principal. Trigger: recurring backlog or a deliverable consuming known time. Why now: existing service assets and local delivery learning (EV02/03), not a claimed sudden model threshold. Loop: select -> baseline -> configure -> comparable executions -> accept/fix/leave unchanged -> handoff. Model: fixed scope, privately priced after qualification. Differentiation: actual delivery, context-specific controls and auditable proof; none is a proven moat. Human/AI architecture: founder implements, assistant processes approved inputs, client decides acceptance. Must be true: ASM01/02/04; acceptable outcomes within a capped delivery budget. Fastest falsification: five qualified conversations with concrete recent examples, followed by a paid bounded pilot. Kill: no paid commitment or review/setup costs erase benefit.

PC02 / long-term software - Workflow Acceptance Ledger. Thesis: teams will pay to maintain accepted workflow versions and revalidate them when inputs/tools change. User: owner of several operational AI workflows, not generic learners. Trigger: second implementation, disputed output, vendor change or owner turnover. Why now: evidence/version problem is plausible (EV06) and process intelligence is an existing category (EV09); small-team gap remains unproven. Loop: record cases -> compare versions -> approve -> observe failures -> retest on meaningful change. Model: workspace subscription only after recurring demand. Differentiation: portable evidence, approved decision rights, repeatable checks; platforms can copy these. Architecture: manual M05 initially, human acceptance M02, optional execution through existing tools. Future automatic scanner and migration require returning to Step 5. Must be true: ASM03/05 and low-burden data capture. Fast falsification: manually maintain records for pilot workflows and offer paid revalidation after a genuine change. Kill: owners do not consult the record, no paid retest, or every workflow needs a custom data model.

PC03 / contrarian - Independent Workflow Acceptance Lab. Thesis: buyers may value evidence about someone else's AI implementation more than coaching or another implementation vendor. User: owner evaluating an existing tool/vendor, not an AI beginner. Trigger: acceptance, renewal, failed rollout or disputed value. Loop: agree rubric -> sample historical cases -> blind replay -> report accept/fix/reject. Model: fixed assessment; repeat only on a real change. Differentiation: evaluator independence and reproducible local evidence; not certification. Architecture: human adjudication with AI-assisted checks, using M02/M05; no large platform required. Must be true: buyers pay separately and provide usable samples. Fast falsification: five buyers with an existing implementation, ask for a paid assessment tied to a live decision. Kill: universally treated as vendor-included QA or access costs exceed assessment economics.

Concept scoring

All 16 dimensions are 1-5, higher favorable. Burden/difficulty/risk dimensions are inverted: 5 means easier/lower burden/safer. Codes: pain P, trigger T, measurable value V, willingness to pay W, frontier resistance F, founder feasibility X, distribution ease D, low human burden H, automation potential A, recurrence R, privacy ease Q, implementation ease I, low platform risk L, defensibility B, first-revenue speed S, five-user evidence E. All scores E0; no W score is measured. I12 means the no-new-product strategy with existing consulting revenue, not paid demand for a nonexistent product.

Baseline weights: 8/5/14/5/14/12/4/4/3/4/3/4/4/3/5/8. Impact/evidence-heavy: 10/5/20/5/15/6/3/2/2/3/2/3/3/2/4/15. Feasibility/risk-heavy: 5/4/10/4/10/18/7/7/2/3/6/7/5/2/5/5. Each sums to 100. Raw inputs and scores appear below and in scores.csv. Load-bearing: no software investment without paid target demand, feasible approved data, and recurring use; current scores cannot override these failures.

ID	Raw P/T/V/W/F/X/D/H/A/R/Q/I/L/B/S/E	Baseline / impact / feasibility	Ranks	Stability
I01	4/4/5/3/3/5/3/3/3/3/5/4/2/5/5	4.00 / 4.12 / 4.02	1/1/1	rank-stable
I02	4/3/4/3/3/3/1/2/4/2/3/4/2/3/3	3.13 / 3.26 / 2.99	8/8/9	rank-stable
I03	4/4/5/2/4/4/2/3/3/3/3/4/4/3/4/5	3.87 / 4.08 / 3.73	2/2/4	rank-stable
I04	3/3/4/2/3/3/2/4/4/4/3/3/2/3/2/3	3.07 / 3.12 / 3.01	9/9/7	rank-stable

ID	Raw P/T/V/W/F/X/D/H/A/R/Q//L/B/S/E	Baseline / impact / feasibility	Ranks	Stability
I05	2/2/3/2/2/2/3/5/3/1/2/1/2/1/2	2.19 / 2.22 / 2.10	14/14/14	rank-stable
I06	4/4/5/3/4/2/2/3/4/4/1/1/3/3/1/3	3.22 / 3.48 / 2.83	7/6/10	rank-unstable
I07	4/4/4/3/3/2/2/2/4/5/1/1/2/3/1/3	2.90 / 3.11 / 2.54	11/10/12	rank-stable
I08	4/5/4/3/3/3/2/2/3/2/2/3/2/4/4	3.23 / 3.44 / 3.01	6/7/8	rank-stable
I09	4/4/5/3/3/4/3/3/4/3/3/4/3/5/5	3.86 / 4.04 / 3.76	3/3/3	rank-stable
I10	4/4/4/3/2/2/2/3/4/4/1/2/2/2/2/3	2.82 / 3.00 / 2.58	12/11/11	rank-stable
I11	2/3/3/1/1/5/3/5/5/1/5/5/1/1/5/4	3.02 / 2.87 / 3.42	10/12/6	rank-unstable
I12	3/4/4/3/3/5/4/3/3/2/4/5/5/2/5/5	3.85 / 3.87 / 4.02	4/4/2	rank-stable
I13	4/3/5/3/4/3/2/2/2/3/4/4/5/2/4/4	3.64 / 3.86 / 3.50	5/5/5	rank-stable
I14	3/3/4/2/4/1/1/3/4/4/1/1/3/4/1/2	2.69 / 2.86 / 2.32	13/13/13	rank-stable

Selection is PC01 for the next commercial test; PC02 is the software option, PC03 the contrasting business. I12's high feasibility is a reason to house PC01 within NY AI Solutions, not evidence the broader opportunity is dead. Close ranks must not be interpreted as meaningful numeric precision.

Business Product Brief: selected PC01

Problem/JTBD: when a repeated task consumes effort or causes rework, help the owner get an acceptable result with less total burden and know whether the change was worth it (P02/P03; RC01/RC02).

Value proposition: 'Get one recurring workflow working better in the tools you already use, with a before/after acceptance record and an operator who can run it.' This is internal draft positioning, not a public ROI promise.

Target: one accessible Israeli professional-services niche; provisional first candidate is consultant/risk-advisory briefing or recurring internal reporting. Founder background in local records suggests fit (EV11), but access must be checked. This is a candidate wedge, not a decision to discard the existing broader niche before the gate.

MVP: one workflow, one owner, approved inputs, one configured route, agreed output rubric, representative cases, operator handoff and short follow-up. Choose existing vertical software if it already meets the requirement.

Non-goals: new app, course, mastery score, passive screen recording, autonomous external actions, marketplace, continuous frontier scanner, or guaranteed percent improvement.

Why this wins provisionally: fits existing delivery infrastructure (EV02), addresses recorded last-mile and measurement gaps (EV03), can obtain evidence sooner than SaaS. Main weakness: competition and zero paid-demand evidence (EV04/EV08).

Economics: measure delivery hours including sales, access setup, evaluation, handoff and support; fee minus tool costs and valued founder time determines contribution. Buyer value is actual capacity/rework benefit, not hours multiplied by billing rate unless that capacity can be monetized. No prices are published by this report.

Metrics: median total human minutes per accepted output, acceptance rate, rework minutes, elapsed cycle time, additional tool cost, unsupported-output incidents, operator success, repeated use, paid commitment. Report sample size and range; five users cannot establish population effect or rare-event safety.

Roadmap: existing fit-call gate -> one paid pilot -> repeated same-family deliveries -> manual ledger -> paid revalidation -> only then consider small software. Each expansion requires fresh evidence, not calendar progression.

Full Product Specification: PC01 experiment, no build authorized

Status: test specification. Personas: budget owner, workflow operator, output reviewer; explicitly identify when they are the same person. Goals: acceptable output, net operational improvement, safe ownership transfer. Non-goals match the brief.

Journey: fit call selects an actual recent workflow -> owner approves data path -> baseline runs logged -> acceptance rubric locked -> live setup booked -> improved process trialed -> comparison reviewed -> accepted or restored -> independent operator run -> follow-up on reuse.

Information architecture: one engagement folder with intake, workflow card, baseline, acceptance rubric, implementation record, run log, comparison, operator instructions and decision. No customer application required.

Objects: Workflow(id, owner, trigger, frequency, inputs, output, constraints); Version(id, workflow_id, tool/config snapshot, change reason); Case(id, redacted input reference, expected properties, difficulty); Run(id, case_id, version_id, active/review/rework minutes, elapsed time, cost, status); Acceptance(id, reviewer, rubric version, findings); DecisionRight(action, permitted actor, approval, fallback); Change(id, reason, decision). Keep client PII outside this repository.

State: candidate -> qualified -> baseline-ready -> configured -> trial -> accepted/fix/rejected -> handed-over -> archived/retest. Reject/fix are valid outcomes. Version changes invalidate prior acceptance for changed behavior until reviewed.

REQ01 (P01/RC03/M03): capture recent occurrence, frequency, owner and economic consequence; reject vague aspirational tasks.

REQ02 (P03/RC02/M02): record baseline before claiming gain; if absent, label before/after unknown.

REQ03 (P02/RC01/M01): live setup and named operator are completion requirements, not optional follow-up.

REQ04 (P03/RC02/M02): predefine acceptance criteria including critical errors; compare cases of similar type/difficulty; record failed and abandoned runs.

REQ05 (P03/RC02/M02): include generation, preparation, review and correction effort; track parallel agent elapsed time separately from human active time.

REQ06 (P04/RC04/M04): manual fallback and escalation owner; customer approves externally consequential actions.

REQ07 (P02/RC01/M04): operator demonstrates one complete run and recovery from a deliberately benign exception without founder intervention.

REQ08 (P03/RC02/M05): keep versions, source references and approval evidence exportable; no model self-grade accepted as final truth.

AI behavior: process only approved minimized content; distinguish extraction from inference; show missing evidence; abstain/escalate on unsupported conclusions; never silently fill missing client facts. Independent rubric plus human review determines acceptance.

Trust/privacy path for initial briefing candidate: model sees approved public sources or redacted excerpts and generic objectives; never client names, contacts, credentials, private case identifiers or distinctive sensitive narratives. Owner reviews redaction because aliases alone may not prevent identification. Identity is reattached locally by the authorized operator after final review when needed. If sensitive context is essential, redesign or defer until an approved handling environment is available; do not claim legal compliance.

UX: begin with the real task, not lessons; show next action, evidence and exceptions; keep technical configuration in operator notes. Clearly label draft vs accepted output and model suggestions vs reviewer decisions.

Edge cases: missing baseline means no uplift claim; insufficient runs means provisional result; absent reviewer pauses acceptance; changed tool behavior triggers retest; denied access stops setup; low task frequency extends observation rather than inventing repetitions; worsened quality restores previous method.

Analytics: minimal manual run log; numerator/denominator for acceptance; medians and ranges; no raw sensitive prompts in central analytics. Record payment/proposal separately from positive feedback. Record second use and second purchase separately.

Nonfunctional requirements: client-controlled storage, least needed access, exportable plain records, reversible configuration, explicit retention agreement and deletion owner. Response-time and availability promises are set only after observing actual tool constraints; no premature SLA.

Acceptance: all REQ01-REQ08 evidenced; no critical error in observed cases; quality meets preset rubric; operator can run and recover; total effort is compared honestly. This is a pilot gate, not a safety certification.

Dependencies/open questions: actual buyer access, approved data, sufficient frequency, available reviewer, tool permissions, founder capacity, pricing and evidence independence. Phase 0 qualification; phase 1 measured service; phase 2 follow-up; software is outside this specification.

Verticalization

Consultants: accessible artifacts and repeated briefs; subjective quality and confidentiality complicate evaluation. Risk consultants: stronger provisional founder fit and costly errors, but sensitive context. Product managers: many varied cross-team tasks, unclear personal budget. Accountants: repeatable workflows, but specialized software and domain controls. Recruiters: repeat volume, sensitive personal data and ATS competition. Marketing agencies: repeated deliverables and rapid feedback, but crowded tools and weak attribution to revenue. Lawyers: high-value work, high domain/review burden. Financial analysts: source traceability possible, but high stakes and specialized incumbents. These are E0 comparative hypotheses, not market rankings established by research.

Choose by reachable buyer + repeated workflow + safe samples + objective acceptance, not occupation label alone. A narrow workflow family is likely more useful initially than 'knowledge workers'; a vertical is only stronger if it improves access/reuse enough to offset domain burden.

Gate and Red Team

CONDITIONAL PASS to strategy. Challenge: this is the existing Value Sprint renamed. Response: retain that offer and change its internal completion standard; do not launch a second brand. Challenge: software option is premature. Response: PC02 remains a manual evidence experiment; M06/M07 excluded. Inputs M01-M05/EV02-EV11; outputs PC01-PC03, REQ01-REQ08.

Source and evidence register

All sources accessed 2026-09-06. Grades apply only to the stated claim, never automatically to this business. E0 = hypothesis; E1 = anecdote/proxy; E2 = working support; E3 = direct primary evidence; E4 = repeated convergent primary evidence. No E4 demand evidence exists.

- SRC01 / EV01: User's Innovation handoff, preserved as source-handoff.txt. E3 for the requested research boundary; E0 for its unverified market and demand assertions. Direction: mixed. Earlier Storm/Red Team full transcripts were not reviewed; their conclusions are represented by this handoff.
- SRC02 / EV02: Local docs/service/offer.md and delivery-workflow.md, dated 2026-08-30. E3 for documented offer design, unknown for willingness to pay. Direction: supports an existing service vehicle. Existing AI Value Sprint: 7-10 business days, one Quick Win implemented.
- SRC03 / EV03: Local docs/clients/2026-08-30-michal-engagement-retro.md. E1 for generalizable delivery learning; E3 for the record's stated limitations. A related-party pilot reached setup, but captured no baseline and does not establish arm's-length demand. Direction: mixed. We do not repeat its claim that it proves the method works as a validated outcome.
- SRC04 / EV04: Local docs/gtm/fit-call-log.md, read today: zero recorded fit calls, qualified workflows, proposal requests and paid engagements. E3 for the record, not proof that no unlogged conversations occurred. Direction: contradicts any claim of commercial validation. Existing review: 2026-09-11 or ten calls.
- SRC05 / EV05: NBER, Generative AI at Work, <https://www.nber.org/papers/w31161> . Primary research surfaced in search; page fetch returned 403. E2 in this run because only the abstract/search evidence was accessible. Direction: supports possible task-specific gains, not universal uplift or consultant demand. Avoid transferring study effect sizes to this offer.
- SRC06 / EV06: METR, 2026-02-24 study-design update, <https://metr.org/blog/2026-02-24-uplift-update/> . E3 for its reported measurement problems; external validity to professional services remains uncertain. Direction: mixed. Newer study selection and timing problems prevent a clean current uplift estimate; the early-2025 slowdown cannot establish current model performance.
- SRC07 / EV07: Anthropic, Building effective agents, <https://www.anthropic.com/engineering/building-effective-agents> . E3 for vendor architecture guidance, E1 for commercial inference. Direction: supports using simple designs and increasing complexity only when justified. Not independent validation of this product.
- SRC08 / EV08: Zapier support and partner documentation, <https://help.zapier.com/hc/en-us/articles/8496213764877-Get-help-and-support-with-Zapier> and <https://docs.zapier.com/integrations/quickstart/app-developer-services> . E3 for availability of implementation/maintenance partners. Direction: contradicts claims that workflow implementation is empty competitive space. No verified competitor prices or small-firm outcomes.

- SRC09 / EV09: Celonis Enterprise AI, <https://www.celonis.com/solutions/ai> . E3 for advertised process intelligence/digital-twin positioning; E0 for efficacy in our target segment. Direction: contradicts novelty of a broad work digital twin. Does not establish a direct small-firm substitute.
- SRC10 / EV10: Microsoft Work Trend Index, <https://www.microsoft.com/en-us/worklab/work-trend-index> . E3 for vendor research agenda; E1 for this opportunity. Direction: supports organizational adaptation as a live topic. Vendor incentives and survey selection apply; not willingness-to-pay evidence.
- SRC11 / EV11: Local docs/service/value-proposition.md. E3 for recorded positioning and founder-background claims, not independent verification. Direction: supports provisional fit with risk-aware professional services; reachable buyers and delivery capacity remain unverified.

Shared assumptions, risks, and contradictions

ASM01: reachable buyers have a costly recurring workflow (E0). ASM02: they will pay for implementation plus proof (E0). ASM03: cases share reusable structure (E0). ASM04: safe representative samples can be obtained (E0). ASM05: customers will later pay for revalidation (E0).

RISK01: verification/setup overhead erases gains. RISK02: no paying buyer. RISK03: bespoke delivery consumes founder capacity. RISK04: sensitive context leaks even after names are removed. RISK05: bundled platform features absorb the software. RISK06: provider grades its own work too generously.

Contradictions: general productivity opportunity (EV05/EV10) coexists with local absence of measured outcomes (EV03/EV04); these concern different populations and require a measured local pilot. More capable agents (scenario, not forecast) can increase delegation while reducing the price of setup; survival of independent assurance remains ASM05. Privacy differentiation in local marketing (EV11) is not proof competitors lack it (EV08).