

AI WORK / INNOVATION

Decision and next-workflow handoff

IDEATION-2026-09-06-AI-WORK-01

IDEATION v0.2 | 6 September 2026

Evidence-led exploration of the Personal AI Work Coach / AI Workflow Transformation opportunity.

Research complete. Commercial evidence gates remain conditional.

Step 8 - Decision and next-workflow handoff

Run: IDEATION-2026-09-06-AI-WORK-01. Version: 1.0. Method: IDEATION v0.2. Owner: Nadav / Codex. Date: 2026-09-06. Status: final desk-research artifact. Boundary: professional AI workflow improvement; no deployment or customer testing. Inputs: source handoff, source register and prior steps in this run.

Decision

DEC01: IDEATION investment decision = CONDITIONAL GO for a bounded commercial validation of PC01 inside the existing NY AI Solutions offer. No software build commitment. Innovation verdict = REFRAME.

Strongest resulting concept: existing AI Value Sprint with a workflow acceptance contract, measured total effort and operator handoff. Strongest new insight: the valuable persistent object may be an accepted workflow version and its evidence, not a course, coach or autonomy score.

What changed from Storm: education and coaching move from product to optional transfer mechanism; the six-week program loses priority. What changed from Red Team: implementation alone is insufficient differentiation; acceptance evidence and ownership are explicit, and the opportunity is reconciled with an offer already operating locally. The long-term software possibility is conditional rather than assumed.

Biggest remaining uncertainty: who pays enough for the incremental implementation/verification work, and whether repeatable delivery leaves adequate contribution after founder effort. EV04 contains zero recorded commercial evidence; EV03 contains no measured outcome. These are load-bearing, not minor caveats.

Strength and weakness pattern

Problem strength: plausible but unmeasured locally. Cause clarity: partial, strongest for setup ownership and missing baselines. Mechanism feasibility: good for a manual bounded service. Differentiation: weak-to-moderate and copyable. Defensibility: unestablished. Business potential: credible enough for a paid test, not proven. Software recurrence: E0. Do not average away demand, data-access or unit-economics failures.

Explicit answers to the 12 questions

1. Learning necessary? Operator judgment and recovery are necessary; lessons and a course are optional.
2. Optimize what? Total human+AI performance, measured through accepted outputs and all associated effort.
3. Correct core object? Sell an outcome, change a workflow, retain an accepted workflow version and evidence.
4. Done-with-you or done-for-you? Done-for-you setup plus done-with-you acceptance and handoff; choose DIY when it suffices.
5. Must successful workflows become agents? No. Delegate only where quality, cost and recoverability justify it.
6. What must humans know? Objectives, input limits, acceptance criteria, decision rights, exception diagnosis, escalation and fallback.
7. Assistance Decay? Track provider rescue decay; embrace permanent productive AI assistance.
8. Strongest retention? Revalidation after a meaningful operational change or failure; separately test willingness to pay.
9. What survives better assistants? Accountable delivery, local evidence and decision ownership may survive. These are economic jobs, not guaranteed platform-resistant features.
10. Software or consulting? Service is the strongest current test. Software is an option conditional on repeated records, reuse and paid maintenance demand.
11. Narrow vertical stronger? Provisionally a narrow workflow family is stronger. Choose a vertical only with access, sample availability and relevant expertise.
12. Regret not testing? A real paid workflow upgrade measured end to end, and a separate buyer's willingness to pay for independent acceptance testing of an existing implementation.

Next workflow: Advisor

Run Advisor next, narrowly, to reconcile PC01 with NY AI Solutions, EXD-026, founder capacity and the approaching 2026-09-11/ten-call gate. Do not create a second venture by default. This is a recommendation; Advisor has not been dispatched and no customer outreach has occurred.

Storm next only if a concrete decision-critical external question would change the selected buyer/workflow, such as whether existing vertical software already solves it. Bound retrieval to that question and demand primary sources. Do not rerun the original broad 'AI learning' research.

Red Team next after a specific buyer, workflow, data path, scope and acceptance protocol exist, and before a consequential implementation commitment. Focus on false ROI, review burden, fragile access, privacy, dependence, liability boundaries and support economics. Another generic concept critique would mostly revisit known objections.

Preferred sequence: Advisor decision -> existing fit-call validation -> focused Red Team on the proposed pilot -> paid measured pilot -> evidence review. Conditional Storm can slot before pilot selection if a factual gap blocks the decision.

Ready-to-paste Advisor handoff

I completed IDEATION-2026-09-06-AI-WORK-01. Treat its verdict as REFRAME and investment decision as CONDITIONAL GO for service validation only. The initial Personal AI Work Coach idea became a workflow acceptance/implementation concept. Our local NY AI Solutions already sells a 7-10 business-day AI Value Sprint with one Quick Win. The current fit-call log has zero recorded calls/proposals/payments; the related-party pilot has setup evidence but no baseline or arm's-length demand. EXD-026 reviews on 2026-09-11 or ten calls. Recheck current records before deciding.

Decide whether this belongs entirely inside the existing practice, which one reachable segment and recurring workflow to test, and how to allocate founder time before the existing gate. Compare the sprint against independent acceptance testing and doing no new product work. Give one recommendation, a bounded next action and a stopping rule. Treat all concept scores as hypotheses; do not restart broad ideation, presume subscription demand, publish prices, promise ROI, build software or contact anyone. If contact access or available hours are unknown, identify the missing decision input. Use the source snapshot, assumptions ASM01-ASM05 and planning thresholds below.

Downstream test brief, recommended not executed

Owner: Nadav. Target workflow: Advisor first, then existing NY AI Solutions fit-call process in this repository. User/outcome hypothesis: an accessible professional-services owner will pay for one accepted workflow improvement with lower total burden. Mechanism: M01/M02/M04 with manual M03/M05. Non-goals: software, large integration project, new brand, generic coaching, agentic promotion by default.

First scope: use five qualified conversations as an early checkpoint within the existing ten-call/date gate. Bring an actual recent recurring task, approved sample and identifiable owner; do not count friendly feedback as paid intent. The report does not reset or override EXD-026.

Pilot scope after buyer agreement: one workflow, 7-10 business days where task frequency permits, plus a follow-up. Suggested cap 12 founder delivery hours for the first bounded workflow; this is a planning assumption, not known feasibility or founder availability. Book setup at kickoff. If adequate observation requires longer, state it rather than compressing evidence.

Planning success thresholds: at least three of five qualified conversations expose a concrete recurring pain; at least one requests a scoped proposal and makes a paid commitment before implementation. Pilot target: at least three baseline and three improved comparable runs, all failures retained; median total human effort at least 20% lower with no material quality decline under a preset rubric; operator completes one independent run and recovery exercise. The 20% target is chosen to test meaningful value, not research-derived or a sales guarantee. Tiny samples are directional; do not claim statistical proof or rare-error safety.

Kill/pause criteria: no paid commitment at the existing gate; no useful safe samples; cost of setup/review exceeds recurring benefit; critical output failure; owner cannot operate within the agreed support cap; or each prospect requires an incompatible offer. Diagnose whether the failure kills the workflow, niche or wider business rather than silently changing thresholds. Existing EXD-026 criteria remain controlling for the practice.

Strong-GO criteria for repeating the service, not SaaS: three independent paid deliveries in one workflow family with accepted outputs, positive measured contribution, and demonstrated reuse. Software exploration requires a recurring evidence-management pain plus at least two buyers agreeing to pay for ongoing revalidation after a genuine change. These

are planning gates and should be reviewed by Advisor.

Instrumentation: source of introduction, qualification, actual paid amount privately, baseline and improved case IDs, total human/review/rework time, elapsed time, tool cost, acceptance decisions, exceptions, founder delivery/support hours, repeat usage and second purchase. Do not record client PII here.

Privacy: approved minimized samples only; owner-controlled storage; redact identifying narrative, not just names; no external sending/production actions without the customer's agreed permissions. Anonymized examples in research; identities reattached by the authorized owner outside model context when necessary.

Cost/time: Advisor should select an actual founder-time bound; proposed decision-session cap 60 minutes and no new software spend. No token budget requested. Proposal/pilot fee and delivery budget require a real scope; no purchase is authorized by this report. Stop at the existing date/call gate and review evidence. User instruction currently authorizes this research and advice only.

What specifically not to build

No standalone AI coach, lesson library, capability graph, frontier-percentage dashboard, universal work digital twin, continuous scanner, agent marketplace or proprietary orchestration platform. No second landing page for an offer already present. Preserve the software hypothesis as a testable option rather than an implementation backlog.

Gate and retrospective

CONDITIONAL PASS for a bounded research decision; paid-market and software gates remain unpassed. All steps completed as desk research, no manual user tests performed. Original scores frozen for later calibration. Main methodological lesson: check the founder's current offer and actual commercial evidence before interpreting a conceptual pivot as a new company. Compare scored feasibility and value with actual delivery hours and paid outcomes after the downstream cycle; do not rewrite predictions retroactively.

Final Innovation Verdict

REFRAME

Source and evidence register

All sources accessed 2026-09-06. Grades apply only to the stated claim, never automatically to this business. E0 = hypothesis; E1 = anecdote/proxy; E2 = working support; E3 = direct primary evidence; E4 = repeated convergent primary evidence. No E4 demand evidence exists.

- SRC01 / EV01: User's Innovation handoff, preserved as source-handoff.txt. E3 for the requested research boundary; E0 for its unverified market and demand assertions. Direction: mixed. Earlier Storm/Red Team full transcripts were not reviewed; their conclusions are represented by this handoff.

- SRC02 / EV02: Local docs/service/offer.md and delivery-workflow.md, dated 2026-08-30. E3 for documented offer design, unknown for willingness to pay. Direction: supports an existing service vehicle. Existing AI Value Sprint: 7-10 business days, one Quick Win implemented.

- SRC03 / EV03: Local docs/clients/2026-08-30-michal-engagement-retro.md. E1 for generalizable delivery learning; E3 for the record's stated limitations. A related-party pilot reached setup, but captured no baseline and does not establish arm's-length demand. Direction: mixed. We do not repeat its claim that it proves the method works as a validated outcome.

- SRC04 / EV04: Local docs/gtm/fit-call-log.md, read today: zero recorded fit calls, qualified workflows, proposal requests and paid engagements. E3 for the record, not proof that no unlogged conversations occurred. Direction: contradicts any claim of commercial validation. Existing review: 2026-09-11 or ten calls.

- SRC05 / EV05: NBER, Generative AI at Work, <https://www.nber.org/papers/w31161> . Primary research surfaced in search; page fetch returned 403. E2 in this run because only the abstract/search evidence was accessible. Direction: supports possible task-specific gains, not universal uplift or consultant demand. Avoid transferring study effect sizes to this offer.

- SRC06 / EV06: METR, 2026-02-24 study-design update, <https://metr.org/blog/2026-02-24-uplift-update/> . E3 for its reported measurement problems; external validity to professional services remains uncertain. Direction: mixed. Newer study selection and timing problems prevent a clean current uplift estimate; the early-2025 slowdown cannot establish

current model performance.

- SRC07 / EV07: Anthropic, Building effective agents, <https://www.anthropic.com/engineering/building-effective-agents> . E3 for vendor architecture guidance, E1 for commercial inference. Direction: supports using simple designs and increasing complexity only when justified. Not independent validation of this product.

- SRC08 / EV08: Zapier support and partner documentation, <https://help.zapier.com/hc/en-us/articles/8496213764877-Get-help-and-support-with-Zapier> and <https://docs.zapier.com/integrations/quickstart/app-developer-services> . E3 for availability of implementation/maintenance partners. Direction: contradicts claims that workflow implementation is empty competitive space. No verified competitor prices or small-firm outcomes.

- SRC09 / EV09: Celonis Enterprise AI, <https://www.celonis.com/solutions/ai> . E3 for advertised process intelligence/digital-twin positioning; E0 for efficacy in our target segment. Direction: contradicts novelty of a broad work digital twin. Does not establish a direct small-firm substitute.

- SRC10 / EV10: Microsoft Work Trend Index, <https://www.microsoft.com/en-us/worklab/work-trend-index> . E3 for vendor research agenda; E1 for this opportunity. Direction: supports organizational adaptation as a live topic. Vendor incentives and survey selection apply; not willingness-to-pay evidence.

- SRC11 / EV11: Local docs/service/value-proposition.md. E3 for recorded positioning and founder-background claims, not independent verification. Direction: supports provisional fit with risk-aware professional services; reachable buyers and delivery capacity remain unverified.

Shared assumptions, risks, and contradictions

ASM01: reachable buyers have a costly recurring workflow (E0). ASM02: they will pay for implementation plus proof (E0). ASM03: cases share reusable structure (E0). ASM04: safe representative samples can be obtained (E0). ASM05: customers will later pay for revalidation (E0).

RISK01: verification/setup overhead erases gains. RISK02: no paying buyer. RISK03: bespoke delivery consumes founder capacity. RISK04: sensitive context leaks even after names are removed. RISK05: bundled platform features absorb the software. RISK06: provider grades its own work too generously.

Contradictions: general productivity opportunity (EV05/EV10) coexists with local absence of measured outcomes (EV03/EV04); these concern different populations and require a measured local pilot. More capable agents (scenario, not forecast) can increase delegation while reducing the price of setup; survival of independent assurance remains ASM05. Privacy differentiation in local marketing (EV11) is not proof competitors lack it (EV08).